

Lecture 09: Queries against several tables


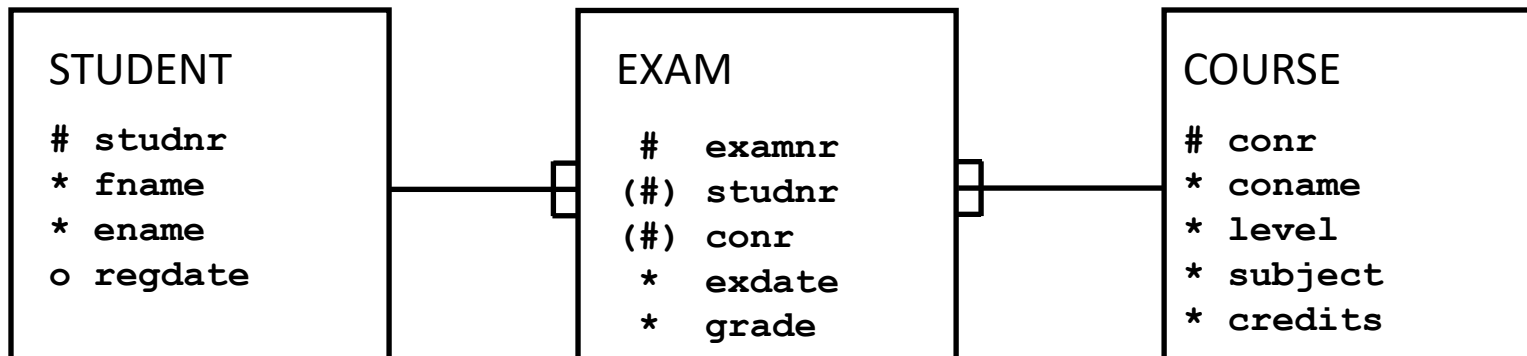


1. *Sub queries (nested queries)*
2. *IN-operator*
3. *Join*
4. *Cartesian product*
5. *Table alias*
6. *Distinct*
7. *Group by*
8. *Having*
9. *In Line View*
10. *Order by*
11. *The pseudo column: rownum*
12. *ANSI Join syntax*

Pär Douhan, pdo@du.se




Exercises





Nested search


- ▶ When columns are shown from only one table.
- ▶ The **IN**-operator is used when the nested query may return more than one row.
- ▶ The sub query must be within (parentheses).
- ▶ The nested queries are executed, in a given order, inside and out.
- ▶ The result is "**distincted**" i.e. no duplicate rows.
- ▶ No sorting (**order by**) in sub queries.





Join

- ▶ Can show columns from more than one table.
- ▶ The result shows with duplicate rows, has to be "**distinct**" in order to remove duplicate rows.






E1

E1: Show STUDNR, FNAME, ENAME for those students who have studied course number 1. Use nested search.

```
select studnr,fname,ename
from student
where studnr in(select studnr
                 from exam
                 where conr = 1);
```



STUDNR	FNAME	ENAME
1	niklas	strömberg
2	arman	dobitnik







E2

E2: Show STUDNR, FNAME, ENAME for those student who have studied a course with the name 'beginning java'. Use sub queries.

```
select studnr, fname, ename
from student
where studnr in (select studnr
                  from exam
                  where conr in (select conr
                                from course
                                where lower(coname) = 'beginning java'));
```



STUDNR	FNAME	ENAME
1	niklas	strömberg
2	arman	dobitnik





E3

E3: Same as E2 but use Equi-Join without table alias.

```
select student.studnr, student.fname, student.ename
from student, exam, course
where lower(course.coname) = 'beginning java'
and student.studnr = exam.studnr
and exam.conr = course.conr;      }  -- Join condition
```

If we forget the join condition, we get cartesian product

STUDNR	FNAME	ENAME
1	niklas	strömberg
2	arman	dobitnik






Cartesian product


A collection is an unordered set of unique data

A database table = a collection.


Collection A



Collection B



The cartesian product of $A \times B$ is the set of all pairs (a,b)



$$A \times B = \{B, X, E\} \times \{5, 2, 8\}$$

$$= \{ (B, 5) , (B, 2) , (B, 8) , (X, 5) , (X, 2) , (X, 8) , (E, 5) , (E, 2) , (E, 8) \}$$



E4

E4: Same as E3 but use table alias.

```
select s.studnr,s.fname,s.ename  
from student s,exam e,course c  
where lower(c.coname) = 'beginning java'  
and s.studnr = e.studnr  
and e.conr = c.conr;
```

Does not work with **as**

STUDNR	FNAME	ENAME
1	niklas	strömberg
2	arman	dobitnik



E5



E5: Show FNAME, ENAME for those students who haven't participated in any course.

```
select initcap(fname||' '||ename) NO_STUDY_RESULT
from student
where studnr not in (select studnr
                     from exam);
```



NO_STUDY_RESULT

Hans Dalros



E6

E6: Show FNAME, ENAME for those students who have participated in a course, order by ename descending. Use join!

```
select s.fname,s.ename
from student s, exam e
where s.studnr = e.studnr  -- Join condition
order by ename desc;
```

FNAME	ENAME
-----	-----
niklas	strömberg
niklas	strömberg
lena	nordin
anna	larsson
jörgen	killling
arman	dobitnik
arman	dobitnik
arman	dobitnik





E7

E7: Same as E6, without duplicates.

```
select distinct s.fname,s.ename
from student s, exam e
where s.studnr = e.studnr
order by enamn desc;
```



FNAME	ENAME
niklas	strömberg
lena	nordin
anna	larsson
jörgen	killin
arman	dobitnik



E8

E8: Same as E7 but use sub queries instead.

```
select fname,ename
from student
where studnr in (select studnr
                  from exam)
order by ename desc;    -- No order by in the sub query!
```

FNAME	ENAME
niklas	strömberg
lena	nordin
anna	larsson
jörgen	killing
arman	dobitnik





E9

E9: Show STUDNR for those students who haven't participated in the course 'AI for Masters'. Join or nested search? Always nesting, with the negation in the outer query.

```
select studnr
from student
where studnr not in(select studnr
                    from exam
                    where conr in(select conr
                                from course
                                where upper(coname) = 'AI FOR MASTERS'));
```

STUDNR

1
2
5
6





E10

E10: Show STUDNR, FNAME, ENAME and the number of courses that the students have participated in.

↓ ↓ ↓

```
select s.studnr,s.fname,s.ename,count(e.examnr) number_of_courses
from student s,exam e
where s.studnr = e.studnr
group by s.studnr,s.fname,s.ename; /* the columns that count() do not work
with! */
```

↑ ↑ ↑

STUDNR	FNAME	ENAME	NUMBER_OF_COURSES
1	niklas	strömberg	2
2	arman	dobitnik	3
3	anna	larsson	1
4	lena	nordin	1
5	jörgen	killing	1





E11

**E11: Same as E10, but also show the students that haven't read any courses.
Outer Join shows NULL-values.**

```
select s.studnr,s.fname,s.ename,nvl(count(e.examnr),0) number_of_courses
from student s,exam e
where s.studnr = e.studnr(+)
group by s.studnr,s.fname,s.ename;
```

STUDNR	FNAME	ENAME	NUMBER_OF_COURSES
1	niklas	strömberg	2
2	arman	dobitnik	3
3	anna	larsson	1
4	lena	nordin	1
5	jörgen	killling	1
6	hans	dalros	0





E12

E12: Show STUDNR, FNAME, ENAME and the total credits that each student has read.

```
select s.studnr,s.fname,s.ename,sum(c.credits) sum_credits
from student s,exam e,course c
where s.studnr = e.studnr
and e.conr = c.conr
group by s.studnr,s.fname,s.ename;
```

STUDNR	FNAME	ENAME	SUM_CREDITS
1	niklas	strömberg	10
2	arman	dobitnik	30
3	anna	larsson	5
4	lena	nordin	5





E13

E13: Same as E12 but sort the result so the student with most credits is on top.

```
select s.studnr,s.fname,s.ename,sum(c.credits) sum_credits
from student s,exam e,course c
where s.studnr = e.studnr
and e.conr = c.conr
group by s.studnr,s.fname,s.ename
order by sum(c.credits) desc;
```

STUDNR	FNAME	ENAME	SUM_CREDITS
2	arman	dobitnik	30
1	niklas	strömberg	10
3	anna	larsson	5
4	lena	nordin	5
5	jörgen	killing	5





E14

E14: Same as E13 but show only the students with 10 credits or more.

```
select s.studnr,s.fname,s.ename,sum(c.credits) sum_credits
from student s,exam e,course c
where s.studnr = c.studnr
and e.conr = c.conr
group by s.studnr,s.fname,s.ename
having sum(c.credits) >= 10 -- condition for group functions
order by sum(c.credits) desc;
```

STUDNR	FNAME	ENAME	SUM_CREDITS
2	arman	dobitnik	30
1	niklas	strömberg	10





E15

E15: Same as E14, but only the top row (with the highest value). Use an 'in line view' to solve the task:

```
select *
from
( -- in line view start
select s.studnr,s.fname,s.ename,sum(c.credits) sum_credits
from student s,exam e,course c
where s.studnr = e.studnr
and e.conr = c.conr
group by s.studnr,s.fname,s.ename
order by sum(c.credit) desc
) -- in line view end
where rownum = 1; -- using Oracle's pseudo column rownum
```



STUDNR	FNAME	ENAME	SUM_CREDITS
2	arman	dobitnik	30



ANSI join syntax

Joins are used to **combine data** from **multiple tables** to form a single result set. Oracle provides **two approaches** to joining tables, the **non-ANSI join** syntax and the **ANSI join** syntax, which look quite different.

The non-ANSI join syntax has historically been the way you perform joins in Oracle and it is still very popular today. The tables to be joined are listed in the FROM clause and the join conditions are defined as predicates in the WHERE clause. Even if you don't like it, you are going to have to get used to it as there is a lot of code out there that still uses it. If you are not familiar with the syntax you will struggle to bug fix any existing code and some of the examples on the internet will look rather mysterious to you.

Many Oracle developers still use the non-ANSI join syntax. Partly this is just because of habit. Partly this is because the Oracle optimizer transforms most ANSI join syntax into the non-ANSI join syntax equivalent before it is executed.



ANSI join syntax

Some join methods are more popular than others, so initially focus your attention on those you are most likely to see. The most common joins you are likely to see in code are the following.

- [INNER] JOIN ... ON
- LEFT [OUTER] JOIN
- RIGHT [OUTER] JOIN


Let's have a look at some examples!




Create tables

```
CREATE TABLE departments (  
  department_id  NUMBER(2) CONSTRAINT departments_pk PRIMARY KEY,  
  department_name VARCHAR2(14),  
  location      VARCHAR2(13));
```

```
CREATE TABLE employees (  
  employee_id  NUMBER(4) CONSTRAINT employees_pk PRIMARY KEY,  
  employee_name VARCHAR2(10),  
  job          VARCHAR2(9),  
  manager_id   NUMBER(4),  
  hiredate     DATE,  
  salary       NUMBER(7,2),  
  commission   NUMBER(7,2),  
  department_id NUMBER(2) CONSTRAINT emp_department_id_fk REFERENCES departments(department_id));
```



[INNER] JOIN ... ON



An INNER JOIN combines data from two tables where there is a match on the joining column(s) in both tables.

Here is an example of an ANSI INNER JOIN:


```
SELECT d.department_name,e.employee_name
FROM   departments d
       JOIN employees e ON d.department_id = e.department_id
WHERE  d.department_id >= 30
ORDER BY d.department_name;
```

DEPARTMENT_NAM	EMPLOYEE_N
SALES	ALLEN
SALES	BLAKE
SALES	JAMES
SALES	MARTIN
SALES	TURNER
SALES	WARD

Here is the non-ANSI equivalent of the previous statement:

```
SELECT d.department_name,e.employee_name
FROM   departments d, employees e
WHERE  d.department_id = e.department_id
AND    d.department_id >= 30
ORDER BY d.department_name;
```


LEFT [OUTER] JOIN



Using the previous example but switching to a LEFT OUTER JOIN means we will see the OPERATIONS department, even though it has no employees.

Here is an example of an ANSI LEFT OUTER JOIN:

```
SELECT d.department_name,e.employee_name
FROM   departments d
       LEFT OUTER JOIN employees e ON d.department_id = e.department_id
WHERE  d.department_id >= 30
ORDER BY d.department_name;
```

DEPARTMENT_NAME	EMPLOYEE_NAME
OPERATIONS	
SALES	ALLEN
SALES	BLAKE
SALES	JAMES
SALES	MARTIN
SALES	TURNER
SALES	WARD


Here is the non-ANSI equivalent of the previous statement:

```
SELECT d.department_name,e.employee_name
FROM   departments d, employees e
WHERE  d.department_id = e.department_id(+)
AND    d.department_id >= 30
ORDER BY d.department_name;
```

Notice the "(+)" is used to indicate the side of the join condition that may be missing. The non-ANSI join syntax is not affected by the order of the tables



RIGHT [OUTER] JOIN



The RIGHT [OUTER] JOIN is the opposite of the LEFT [OUTER] JOIN. It returns all valid rows from the table on the right side of the JOIN keyword, along with the values from the table on the left side, or NULLs if a matching row doesn't exist. All points raised in the previous section apply here also.

The following example has altered the order of the tables so a RIGHT [OUTER] JOIN is now required:

DEPARTMENT_NAME	EMPLOYEE_NAME

OPERATIONS	
SALES	ALLEN
SALES	BLAKE
SALES	JAMES
SALES	MARTIN
SALES	TURNER
SALES	WARD

```
SELECT d.department_name,  
       e.employee_name  
FROM   employees e  
       RIGHT OUTER JOIN departments d ON e.department_id = d.department_id  
WHERE  d.department_id >= 30  
ORDER BY d.department_name, e.employee_name;
```

Remember, the non-ANSI outer join syntax is not dependent on table order, so there is no real concept of right or left outer joins, just outer joins.



The End

